Task:

The MNIST handwritten digit dataset is one of the most commonly used datasets for image classification.

You can access the dataset here: http://yann.lecun.com/exdb/mnist/.

It is also available in many ML packages such as TensorFlow, PyTorch, FastAI. Feel free to directly load the dataset via either source and train a model on it (high accuracy is NOT a criteria for evaluation).

Artifacts needed for evaluation:

1. **Entire codebase with train and inference script to GitHub with links to access**
2. **Saved exported file (use any model repo framework or google drive)**
3. **Dockerfile to build docker image to launch the inference as a REST endpoint (in the same git)**
4. **Screenshot of the post request done locally using postman or curl**

Criteria:

1. **Low latency for inference**
2. **Low inference docker image size**
3. **Code structure**

Bonus Points:

1. **YAML file related to deployment on Kubernetes**
2. **Use of frameworks like Kserve/Triton/MLFlow/Prometheus**
3. **Architecture diagram or writeup explaining flow of requests**